



# Red Hat

## Red Hat Developer Hub 1.8

### Accelerate AI development with OpenShift AI Connector for Red Hat Developer Hub

Installing, configuring, and troubleshooting OpenShift AI Connector for Red Hat Developer Hub



# Red Hat Developer Hub 1.8 Accelerate AI development with OpenShift AI Connector for Red Hat Developer Hub

---

Installing, configuring, and troubleshooting OpenShift AI Connector for Red Hat Developer Hub

## Legal Notice

Copyright © Red Hat.

The text of and illustrations in this document are licensed by Red Hat under a Creative Commons Attribution–Share Alike 3.0 Unported license ("CC-BY-SA"). An explanation of CC-BY-SA is available at

<http://creativecommons.org/licenses/by-sa/3.0/>

. In accordance with CC-BY-SA, if you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, JBoss, OpenShift, Fedora, the Infinity logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux® is the registered trademark of Linus Torvalds in the United States and other countries.

Java® is a registered trademark of Oracle and/or its affiliates.

XFS® is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL® is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js® is an official trademark of Joyent. Red Hat Software Collections is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack® Word Mark and OpenStack logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation, or the OpenStack community.

All other trademarks are the property of their respective owners.

## Abstract

As a developer, when you require access to centralized AI/ML services, you can integrate AI models and model servers from Red Hat OpenShift AI directly into the Red Hat Developer Hub (RHDH) Catalog, so that you can provide a single, consistent hub for discovering, managing, and consuming all components, accelerating time-to-market.

## Table of Contents

<b>CHAPTER 1. UNDERSTAND HOW AI ASSETS MAP TO THE RED HAT DEVELOPER HUB CATALOG .....</b>	<b>3</b>
1.1. MODEL-TO-ENTITY MAPPING	3
1.2. OUT-OF-THE-BOX ASSET DETAILS SYNCED FROM RHOAI	4
<b>CHAPTER 2. SETTING UP OPENSHIFT AI CONNECTOR FOR RED HAT DEVELOPER HUB WITH RED HAT OPENSHIFT AI .....</b>	<b>5</b>
<b>CHAPTER 3. ENRICH AI MODEL METADATA FOR ENHANCED RED HAT DEVELOPER HUB EXPERIENCE</b>	<b>11</b>
3.1. POPULATING THE API DEFINITION TAB IN RHDH API ENTITIES	12
<b>CHAPTER 4. TROUBLESHOOTING CONNECTOR FUNCTIONALITY .....</b>	<b>14</b>
4.1. CHECKING DYNAMIC PLUGINS STATUS	14
4.2. INSPECTING PLUGIN LOGS	14
4.3. INSPECTING THE OPENSHIFT AI CONNECTOR FOR RHDH	14
4.4. OPENSHIFT AI MODEL REGISTRY AND MODEL CATALOG QUERIES	15



# CHAPTER 1. UNDERSTAND HOW AI ASSETS MAP TO THE RED HAT DEVELOPER HUB CATALOG



## IMPORTANT

This section describes Developer Preview features in the OpenShift AI Connector for Red Hat Developer Hub plugin. Developer Preview features are not supported by Red Hat in any way and are not functionally complete or production-ready. Do not use Developer Preview features for production or business-critical workloads. Developer Preview features provide early access to functionality in advance of possible inclusion in a Red Hat product offering. Customers can use these features to test functionality and provide feedback during the development process. Developer Preview features might not have any documentation, are subject to change or removal at any time, and have received limited testing. Red Hat might provide ways to submit feedback on Developer Preview features without an associated SLA.

For more information about the support scope of Red Hat Developer Preview features, see [Developer Preview Support Scope](#).

The OpenShift AI Connector for Red Hat Developer Hub (OpenShift AI Connector for RHDH) serves as a crucial link, enabling the discovery and accessibility of AI assets managed within the Red Hat OpenShift AI offering directly within your RHDH instance.

For more information on model registry components, see [Overview of model registries and model catalog](#).

## 1.1. MODEL-TO-ENTITY MAPPING

Model-to-Entity mapping integrates with OpenShift AI Connector for RHDH, the model catalog, and KServe-based Model Deployments (InferenceServices). This integration automatically converts your AI/ML artifacts into familiar Backstage entities, simplifying management and providing a unified view of your available AI models to your developer teams.

This offering interfaces with the OpenShift AI Connector for RHDH, model catalog, and KServe-based Model Deployments (InferenceServices) to create familiar Backstage entities.

RHOAI Artifact	RHDH/Backstage Entity Kind	RHDH/Backstage Entity Type	Purpose
Model Server (InferenceService)	Component	<b>model-server</b>	Represents a running, accessible AI model endpoint. See <a href="#">Configuring your model-serving platform</a> .
AI Model (Model Registry Version)	Resource	<b>ai-model</b>	Represents the specific AI model artifact, for example, <b>Llama-3-8B</b> .

RHOAI Artifact	RHDH/Backstage Entity Kind	RHDH/Backstage Entity Type	Purpose
Model Server API Details	API	<b>openapi</b> (Default)	Provides the OpenAPI/Swagger specification for the REST endpoint of the model. See <a href="#">Red Hat OpenShift AI: API Tiers</a>
Model Cards	TechDocs	N/A	Model cards from the RHOAI model catalog are associated with the Component and Resource entities. See <a href="#">Registering a model from the model catalog</a> .

Once the OpenShift AI Connector for RHDH is installed and connected with RHOAI, the transfer of information commences automatically.

## 1.2. OUT-OF-THE-BOX ASSET DETAILS SYNCED FROM RHOAI

The connector propagates the following key data:

- InferenceServices (Component type model-server):
  - URL of the OpenShift Route (if exposed).
  - URL of the Kubernetes Service.
  - Authentication requirement status.
- Model registry (Resource type **ai-model**):
  - Model description, artifact URLs, and author/owner information.
- Model catalog:
  - Links to the Model Card (as RHDH TechDocs).
  - Model license URL.

# CHAPTER 2. SETTING UP OPENSHIFT AI CONNECTOR FOR RED HAT DEVELOPER HUB WITH RED HAT OPENSHIFT AI

The installation of the OpenShift AI Connector for Red Hat Developer Hub requires manual updates to RHDH-related Kubernetes resources.

## RHOAI Prerequisites

- To import model cards from the model catalog into TechDocs, you must use RHOAI 2.25.



### NOTE

If you upgraded to RHOAI 2.25 from an earlier version, you must manually enable the model catalog dashboard and model registry before you can import model cards.

- If you used the model catalog in earlier versions of RHOAI, TechDocs propagation does not work for any models you registered into the model registry while at those earlier versions; only models registered into model registry from a RHOAI 2.25 model catalog have their model cards transferred to RHDH as TechDocs.
- For the rest of the features, version 2.20 or later suffices. Enabling model registry and its associated dashboard allows for a user experience that more directly allows for customizing AI Model metadata. For best overall experience, RHOAI 2.25 is recommended.

For more details, see [Enabling the model registry component](#).

## Procedure

1. Configure RHOAI-related RBAC and credentials. A Kubernetes **ServiceAccount** and a **service-account-token** Secret are required for the connector to retrieve data from RHOAI. The following resources must be created, replacing namespace names (**ai-rhdh** for RHDH, **rhoai-model-registries** for RHOAI) as needed:

- **ServiceAccount (rhdh-rhoai-connector)**. For example:

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: rhdh-rhoai-connector
  namespace: ai-rhdh
```

- **ClusterRole and ClusterRoleBinding (rhdh-rhoai-connector)** to allow access to OCP resources like **routes**, **services**, and **inferenceservices**. For example:

```
# Example for `ClusterRole`
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: rhdh-rhoai-connector
  annotations:
    argocd.argoproj.io/sync-wave: "0"
rules:
```

```

- apiGroups:
  - apiextensions.k8s.io
resources:
- customresourcedefinitions
verbs:
- get
- apiGroups:
  - route.openshift.io
resources:
- routes
verbs:
- get
- list
- watch
- apiGroups: [""]
resources:
- serviceaccounts
- services
verbs:
- get
- list
- watch

- apiGroups: ["serving.kserve.io"]
resources: ["inferenceservices"]
verbs: ["get", "list", "watch"]

```

```

# Example for `ClusterRoleBinding`
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: rhdh-rhoai-connector
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: rhdh-rhoai-connector
subjects:
- kind: ServiceAccount
  name: rhdh-rhoai-connector
  namespace: ai-rhdh

```

- **Role** and **RoleBinding** to allow ConfigMap updates within the RHDH namespace (**ai-rhdh**).  
For example:

```

# Example for `Role` and `Rolebinding` in the {product-very-short} namespace ('ai-rhdh')
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: rhdh-rhoai-connector
  namespace: ai-rhdh
rules:
- apiGroups: [""]
  resources: ["configmaps"]
  verbs: ["get", "list", "watch", "create", "update", "patch"]
---
```

```

apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: rhdh-rhoai-connector
  namespace: ai-rhdh
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: Role
  name: rhdh-rhoai-connector
subjects:
  - kind: ServiceAccount
    name: rhdh-rhoai-connector
    namespace: ai-rhdh

```

- **RoleBinding** in the RHOAI namespace (**rhoai-model-registries**) to grant the RHDH **ServiceAccount** read permissions to the model registry data (binding to **registry-user-modelregistry-public**).

```

# Example for `RoleBinding` in the {rhoai-short} namespace (rhoai-model-registries)
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  # if using ODH then change rhoai to odh in the name and namespace here
  name: rhdh-rhoai-dashboard-permissions
  # namespace: odh-model-registries
  namespace: rhoai-model-registries
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: Role
  name: registry-user-modelregistry-public
subjects:
  - apiGroup: rbac.authorization.k8s.io
    kind: Group
    name: system:serviceaccounts:ai-rhdh

```

- Secret (**rhdh-rhoai-connector-token**) of type **kubernetes.io/service-account-token** that goes along with the **rhdh-rhoai-connector ServiceAccount**.

```

apiVersion: v1
kind: Secret
metadata:
  name: rhdh-rhoai-connector-token
  namespace: ai-rhdh
  annotations:
    kubernetes.io/service-account.name: rhdh-rhoai-connector
type: kubernetes.io/service-account-token

```

## 2. Update your RHDH dynamic plugin configuration. The RHDH Pod requires two dynamic plugins.

- In your RHDH dynamic plugins ConfigMap, add the following code:

```

plugins:
  - disabled: false
    package: oci://ghcr.io/redhat-developer/rhdh-plugin-export-overlays/red-hat-developer-
hub-backstage-plugin-catalog-backend-module-model-catalog:bs_1.42.5__0.7.0!red-hat-

```

```

    developer-hub-backstage-plugin-catalog-backend-module-model-catalog
      - disabled: false
        package: oci://ghcr.io/redhat-developer/rhdh-plugin-export-overlays/red-hat-developer-
        hub-backstage-plugin-catalog-techdoc-url-reader-backend:bs_1.42.5__0.3.0!red-hat-
        developer-hub-backstage-plugin-catalog-techdoc-url-reader-backend

```

3. Add the **Connector** sidecar containers to the RHDH Pod.

- If RHDH was installed using the Operator, modify your RHDH custom resource (CR) instance.
- If RHDH was installed using the Helm charts, modify the **Deployment** specification.

4. The system relies on three sidecar containers (OpenShift AI Connector for Red Hat Developer Hub) running alongside the **backstage-backend** container.

Add these sidecar containers to your configuration referencing the **rhdh-rhoai-connector-token**

Secret: **location:** Provides the REST API for RHDH plugins to fetch model metadata **storage-rest:**

Maintains a cache of AI Model metadata in a ConfigMap called **bac-import-model**. \*\* **rhoai-normalizer:** Acts as a Kubernetes controller and RHOAI client, normalizing RHOAI metadata for the connector. The following code block is an example:

+

```

spec:
template:
  spec:
    containers:
      - env:
          - name: NORMALIZER_FORMAT
            value: JSONArrayFormat
          - name: POD_IP
            valueFrom:
              fieldRef:
                fieldPath: status.podIP
          - name: POD_NAMESPACE
            valueFrom:
              fieldRef:
                fieldPath: metadata.namespace
    envFrom:
      - secretRef:
          name: rhdh-rhoai-connector-token
    image: quay.io/redhat-ai-dev/model-catalog-location-
service@sha256:c4471e07be6e0dbe821613053e6264a552cacda7f8604dbf306e6ac9e81e8ab9
    imagePullPolicy: Always
    name: location
    ports:
      - containerPort: 9090
        name: location
        protocol: TCP
    volumeMounts:
      - mountPath: /opt/app-root/src/dynamic-plugins-root
        name: dynamic-plugins-root
    workingDir: /opt/app-root/src
  - env:
      - name: NORMALIZER_FORMAT

```

```

        value: JSONArrayFormat
    - name: STORAGE_TYPE
        value: ConfigMap
    - name: BRIDGE_URL
        value: http://localhost:9090
    - name: POD_IP
        valueFrom:
            fieldRef:
                fieldPath: status.podIP
    - name: POD_NAMESPACE
        valueFrom:
            fieldRef:
                fieldPath: metadata.namespace
    envFrom:
        - secretRef:
            name: rhdh-rhoai-connector-token
    image: quay.io/redhat-ai-dev/model-catalog-storage-
rest@sha256:398095e7469e86d84b1196371286363f4b7668aa3e26370b4d78cb8d4ace1dc9
        imagePullPolicy: Always
    name: storage-rest
    volumeMounts:
        - mountPath: /opt/app-root/src/dynamic-plugins-root
            name: dynamic-plugins-root
    workingDir: /opt/app-root/src
- env:
    - name: NORMALIZER_FORMAT
        value: JSONArrayFormat
    - name: POD_IP
        valueFrom:
            fieldRef:
                fieldPath: status.podIP
    - name: POD_NAMESPACE
        valueFrom:
            fieldRef:
                fieldPath: metadata.namespace
    envFrom:
        - secretRef:
            name: rhdh-rhoai-connector-token
    image: quay.io/redhat-ai-dev/model-catalog-rhoai-
normalizer@sha256:9f19742450a3a9c6d9c01d8341a20db7eb5a52a39348f488ae06b6aa49754a26
        imagePullPolicy: Always
    name: rhoai-normalizer
    volumeMounts:
        - mountPath: /opt/app-root/src/dynamic-plugins-root
            name: dynamic-plugins-root
    workingDir: /opt/app-root/src
    args:
        - '--metrics-address=:8081'

```

1. Enable **Connector** in your **RHDHapp-config.yaml** file. In your **Backstage `app-config.extra.yaml`** file, configure **Entity Provider** under the **catalog.providers** section:

```

    providers:
        modelCatalog:
            development:
                baseUrl: http://localhost:9090

```

where:

**modelCatalog**

Specifies the name of the provider.

**development**

Defines future connector capability beyond a single **baseUrl**.

**baseUrl**

For Developer Preview, this value is the only one supported. Future releases might support external routes.

## CHAPTER 3. ENRICH AI MODEL METADATA FOR ENHANCED RED HAT DEVELOPER HUB EXPERIENCE

While RHOAI provides essential data, an AI platform engineer using RHOAI can enrich the Backstage/RDH experience by adding **custom properties** to the **ModelVersion** or **RegisteredModel** (or annotations to the **KServe InferenceService** if the model registry is not used) so that the OpenShift AI Connector for Red Hat Developer Hub can add the information to the RHDH entities it creates. For more details, see [Editing model version metadata in a model registry](#).

Property Key	Entity Field Populated	Description
<b>API Spec</b>	API Definition Tab	The OpenAPI / Swagger JSON specification for the model REST API.
<b>API Type</b>	API Type	Correlates to supported RHDH/Backstage API types (defaults to <b>openapi</b> ).
<b>TechDocs</b>	TechDocs	URL pointing to a Git repository that follows RHDH TechDocs conventions for the Model Card. Use this setting only if the <b>Model Card to TechDocs</b> mapping is not active.
<b>Homepage URL</b>	Links	A URL considered the home page for the model.
<b>Owner</b>	Owner	Overrides the default OpenShift user as the entity owner.
<b>Lifecycle</b>	Lifecycle	Serves a means to express the lifecycle notion of RHDH/Backstage.
<b>How to use</b>	Links	A URL that points to usage documentation.
<b>License</b>	Links	A URL to the license file of the model.

Property Key	Entity Field Populated	Description
<b>rhdh.modelcatalog.io/model-name</b>	Annotations	A key piece of metadata is the name of the model used when communicating with the model server's REST API. OpenShift AI Connector for Red Hat Developer Hub stores this name in the <b>rhdh.modelcatalog.io/model-name</b> annotation on the <b>Resource</b> entity, and will default this annotation's value to the combined names of the <b>ResourceModel</b> and <b>ModelVersion</b> in the model registry, or the <b>KServe InferenceService</b> name if the model registry is not used, as those often line up with the model name used when communicating with the model server's REST API. But the names are not guaranteed to match. When the names do not match, provide the correct model name for model REST API invocations.

### 3.1. POPULATING THE API DEFINITION TAB IN RHDH API ENTITIES

Because RHOAI does not expose the OpenAPI specification by default, the AI platform engineer can take the following steps to provide this information:

#### Procedure

1. Retrieve OpenAPI JSON: Use a tool such as **curl** to fetch the specification directly from the running endpoint of the AI model server. The following command provides the precise endpoint (**/openapi.json**) and shows how to include a **Bearer** token if the model requires authentication for access.

```
$ curl -k -H "Authorization: Bearer $MODEL_API_KEY"
https://$MODEL_ROOT_URL_INCLUDING_PORT/openapi.json | jq > open-api.json
```

2. Set Property in RHOAI.

- a. In the **RHOAI** dashboard, go to **Model Registry** and select the appropriate **Model Version**.



#### NOTE

We recommend using **Model Version** instead of **Registered Model** to maintain stability if the API changes between versions.

- b. In the **Properties** section, set a key/value pair where the key is **API Spec** and the value is the entire JSON content from the **open-api.json** file.
3. Propagation: The OpenShift AI Connector for Red Hat Developer Hub periodically polls the RHOAI Model Registry, propagates this JSON, and renders the interactive API documentation in the **Definition** tab of the RHDH API entity.

## CHAPTER 4. TROUBLESHOOTING CONNECTOR FUNCTIONALITY

The connector system consists of the two dynamic plugins and the three OpenShift AI Connector for RHDH sidecar containers. Generally speaking, the logs collected must be provided to Red Hat Support for analysis.

The actual contents of the diagnostic data are not part of any product guaranteed specification, and can change at any time.



### NOTE

During startup, you may see non-critical log errors, such as **in cluster config error: open /var/run/secrets/kubernetes.io/serviceaccount/token: no such file or directory**, in the sidecar logs. This error is expected during the initial setup and do not indicate a failure, provided the container eventually becomes healthy.

### 4.1. CHECKING DYNAMIC PLUGINS STATUS

Validate that the dynamic plugins have been successfully installed into your RHDH project Pod by using the following command:

```
$ oc logs -c install-dynamic-plugins deployment/<your RHDH deployment>
```

In the **install-dynamic-plugin** logs, you can check the following installation logs for successful logs:

- **red-hat-developer-hub-backstage-plugin-catalog-backend-module-model-catalog** (Entity Provider)
- **red-hat-developer-hub-backstage-plugin-catalog-techdoc-url-reader-backend** (TechDoc URL Reader)

### 4.2. INSPECTING PLUGIN LOGS

View the OpenShift AI Connector for Red Hat Developer Hub plugins in the **backstage-backend** container. Items to look for:

Plugin Component	Logger Service Target	Common Log Text
Model Catalog Entity Provider	<b>ModelCatalogResourceEntityProvider</b>	<b>Discovering ResourceEntities from Model Server...</b>
Model Catalog TechDoc URL Reader	<b>ModelCatalogBridgeTechdocUrlReader</b>	<b>ModelCatalogBridgeTechdocUrlReader.readUrl</b>

To enable debug logging, set the **LOG\_LEVEL** environment variable to **debug** on the **backstage-backend** container. For more information, see [Monitoring and logging](#).

### 4.3. INSPECTING THE OPENSHIFT AI CONNECTOR FOR RHDH

The OpenShift AI Connector for RHDH sidecars manage the data fetching and storage:

1. Check Cached Data (ConfigMap): The processed AI Model metadata is stored in a **ConfigMap**.

```
$ oc get configmap bac-import-model -o json | jq -r '.binaryData | to_entries[] | "==== \(.key)  
====\n" + (.value | @base64d | fromjson | .body | @base64d | fromjson | tostring)' | jq -R 'if  
startswith("==== ") then . else (. | fromjson) end'
```

2. Check Location Service API: Confirm the location service is providing data to the RHDH Entity Provider.

```
$ oc rsh -c backstage-backend deployment/<your RHDH deployment>  
$ curl http://localhost:9090/list
```

3. Check Sidecar Container Logs:

```
$ oc logs -c rhoai-normalizer deployment/<your {product-very-short} deployment>  
$ oc logs -c storage-rest deployment/<your {product-very-short} deployment>  
$ oc logs -c location deployment/<your {product-very-short} deployment>
```

## 4.4. OPENSHIFT AI MODEL REGISTRY AND MODEL CATALOG QUERIES

To access the same RHOAI data as the connector, use **curl** to query the RHOAI model registry and model catalog APIs, ensuring the **ServiceAccount** token has correct access control:

- Example showing how to fetch registered models

```
$ curl -k -H "Authorization: Bearer $TOKEN"  
$RHOAI_MODEL_REGISTRY_URL/api/model_registry/v1alpha3/registered_models | jq
```

- Example showing how to fetch model versions

```
$ curl -k -H "Authorization: Bearer $TOKEN"  
$RHOAI_MODEL_REGISTRY_URL/api/model_registry/v1alpha3/model_versions | jq
```

- Example showing how to fetch model artifacts

```
$ curl -k -H "Authorization: Bearer $TOKEN"  
$RHOAI_MODEL_REGISTRY_URL/api/model_registry/v1alpha3/model_artifacts | jq
```

- Example showing how to fetch inference services

```
$ curl -k -H "Authorization: Bearer $TOKEN"  
$RHOAI_MODEL_REGISTRY_URL/api/model_registry/v1alpha3/inference_services | jq
```

- Example showing how to fetch serving environments

```
$ curl -k -H "Authorization: Bearer $TOKEN"  
$RHOAI_MODEL_REGISTRY_URL/api/model_registry/v1alpha3/serving_environments | jq
```

- Example showing how to fetch catalog sources

```
$ curl -k -H "Authorization: Bearer $TOKEN"  
$RHOAI_MODEL_CATALOG_URL/api/model_catalog/v1alpha1/sources | jq
```